Centenary Celebrated Sharnbasveshwar Vidya Vardhaka Sangha's











A State Private University approved by Govt. of Karnataka vide Notification No. ED 144 URC 2016 dated 29-07-2017 Recognised by UGC under Section 2f vide No. F.8-29/2017 (CPP-I/PU), dated 20-12-2017 & AICTE, CoA, PCI New Delhi

DEPARTMENT OF ARTIFICIAL INTELLIGENCE & DATA SCIENCE

FACULTY OF ENGINEERING AND TECHNOLOGY(CO-ED)

SHARNBASVA UNIVERSITY KALABURAGI



BIG DATA ANALYTICS

Lab Manual

For 6th Semester

Course Code: 22ADL67

By

Dr. Laxmi Math

List of Experiments

Sl. NO	Title of the Experiment
1	Install Hadoop and Implement the following file management tasks in Hadoop: Adding files and directories,Retrieving files,Deleting files and directories. Hint: A typical Hadoop workflow creates data files (such as log files) elsewhere and copies them into HDFS using one of the above command line utilities.
2	Develop a MapReduce program to implement Matrix Multiplication
3	Develop a Map Reduce program that mines weather data and displays appropriate messages indicating the weather conditions of the day.
4	Develop a MapReduce program to find the tags associated with each movie by analyzing movie lens data.
5	Implement Functions: Count – Sort – Limit – Skip – Aggregate using MongoDB
6	Develop Pig Latin scripts to sort, group, join, project, and filter the data.
7	Use Hive to create, alter, and drop databases, tables, views, functions, and indexes.
8	Implement a word count program in Hadoop and Spark.
9	Use CDH (Cloudera Distribution for Hadoop) and HUE (Hadoop User Interface) to analyze data and generate reports for sample datasets

1. Install Hadoop and Implement the following file management tasks in Hadoop: Adding files and directories, Retrieving files, Deleting files and directories.

Hint: A typical Hadoop workflow creates data files (such as log files) elsewhere and copies them into HDFS using one of the above command line utilities.

Lab: Installing Hadoop and Implementing File Management Tasks in Ubuntu

This lab will guide you through installing Hadoop on Ubuntu and performing basic file management tasks using Hadoop Distributed File System (HDFS), such as adding files and directories, retrieving files, and deleting files and directories.

2. Step-by-Step Guide

Step 1: Install Java

Hadoop requires Java to run. You can install Java 8 or 11 by running the following commands:

sudo apt update

sudo apt install openjdk-8-jdk

Verify Java installation:

java -version

The output should show Java 8 or 11 version, depending on which you installed.

Step 2: Download and Install Hadoop

Download the latest stable version of Hadoop from the official Apache Hadoop website.

wget https://downloads.apache.org/hadoop/common/stable/hadoop-3.3.3.tar.gz

Extract the downloaded file:

tar -xzvf hadoop-3.3.3.tar.gz

Move the extracted files to the /usr/local directory:

sudo mv hadoop-3.3.3 /usr/local/hadoop

Set up environment variables to configure Hadoop:

Open the .bashrc file in a text editor:

nano ~/.bashrc

Add the following lines at the end of the file:

export HADOOP_HOME=/usr/local/hadoopexport
PATH=\$PATH:\$HADOOP_HOME/bin:\$HADOOP_HOME/sbinexport
HADOOP_CONF_DIR=\$HADOOP_HOME/etc/hadoopexport

HADOOP_MAPRED_HOME=\$HADOOP_HOMEexport HADOOP_COMMON_HOME=\$HADOOP_HOMEexport HADOOP_HDFS_HOME=\$HADOOP_HOMEexport YARN_HOME=\$HADOOP_HOMEexport JAVA_HOME=/usr/lib/jvm/java-8-openjdk-amd64

Source the .bashrc to apply the changes:

source ~/.bashrc

Verify Hadoop is installed correctly:

hadoop version

Step 3: Configure Hadoop

Open the Hadoop configuration directory:

cd /usr/local/hadoop/etc/hadoop

Edit core-site.xml to configure the Hadoop file system (HDFS) settings:

nano core-site.xml

Add the following configuration inside the <configuration> tags:

xml

CopyEdit

cproperty>

<name>fs.defaultFS</name>

<value>hdfs://localhost:9000</value></property>

Edit hdfs-site.xml to configure HDFS storage:

nano hdfs-site.xml

Add the following configuration inside the <configuration> tags:

xml

CopyEdit

cproperty>

<name>dfs.replication</name>

<value>1</value></property>

cproperty>

```
<name>dfs.namenode.name.dir</name>
  <value>file:///usr/local/hadoop/hdfs/name</value></property>
cproperty>
  <name>dfs.datanode.data.dir</name>
  <value>file:///usr/local/hadoop/hdfs/data</value></property>
Edit mapred-site.xml to configure MapReduce:
cp mapred-site.xml.template mapred-site.xml
nano mapred-site.xml
Add the following configuration:
xml
CopyEdit
cproperty>
  <name>mapreduce.framework.name</name>
  <value>yarn</value></property>
Edit yarn-site.xml to configure YARN:
nano yarn-site.xml
Add the following configuration:
xml
CopyEdit
cproperty>
  <name>yarn.resourcemanager.hostname</name>
  <value>localhost</value></property>
```

Step 4: Format the Hadoop NameNode

Before starting Hadoop, you need to format the Hadoop NameNode:

hdfs namenode -format

Step 5: Start Hadoop Daemons

Start the Hadoop daemons for HDFS and YARN:

Start HDFS (NameNode and DataNode):

start-dfs.sh

Start YARN (ResourceManager and NodeManager):

start-yarn.sh

Check if all daemons are running:

jps

You should see processes like NameNode, DataNode, ResourceManager, and NodeManager.

6. HDFS File Management Tasks

Now that Hadoop is set up, let's perform basic file management tasks using HDFS.

Task 1: Adding Files and Directories to HDFS

- 1.Create a directory in HDFS:
- 2.hadoop fs -mkdir /user/hadoop/input
- 1.Copy a local file to HDFS:
- 2.hadoop fs -put localfile.txt /user/hadoop/input/

Task 2: Retrieving Files from HDFS

Copy a file from HDFS to the local file system:

hadoop fs -get /user/hadoop/input/localfile.txt /home/ubuntu/

Task 3: Deleting Files and Directories in HDFS

Delete a file in HDFS:

hadoop fs -rm /user/hadoop/input/localfile.txt

Delete a directory recursively:

hadoop fs -rm -r /user/hadoop/input/

7. Verifying File Management Tasks

To list files in a directory in HDFS:

hadoop fs -ls /user/hadoop/input/

To check if a file exists:

hadoop fs -test -e /user/hadoop/input/localfile.txt

If the file exists, it will return 0, else it will return 1.

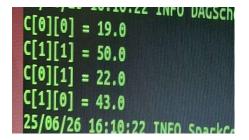
8. Stopping Hadoop Daemons		
After completing your tasks, stop the Hadoop daemons:		
Stop HDFS:		
stop-dfs.sh		
Stop YARN:		
stop-yarn.sh		
DEPARTMENT OF ARTIFICIAL INTELLIGENCE & DATA SCIENCE BIG DATA ANALYTICS-Lab Manual		

2. Develop a MapReduce program to implement Matrix Multiplication

```
from pyspark import SparkContext
sc = SparkContext("local", "MatrixMultiply")
# Read files
linesA = sc.textFile("matrix a.txt")
linesB = sc.textFile("matrix b.txt")
# Parse input
def parse matrix A(line):
  _, i, j, value = line.split(',')
  return (int(i), int(j), float(value))
def parse matrix B(line):
  _, j, k, value = line.split(',')
  return (int(j), int(k), float(value))
matrixA = linesA.map(parse matrix A)
matrixB = linesB.map(parse_matrix_B)
# Emit for multiplication: key = i (shared dimension)
\# A: (i, j, a \ ij) => (j, ('A', i, a \ ij))
# B: (j, k, b_jk) => (j, ('B', k, b_jk))
mappedA = matrixA.map(lambda x: (x[1], ('A', x[0], x[2])))
mappedB = matrixB.map(lambda x: (x[0], ('B', x[1], x[2])))
# Join on j
joined = mappedA.union(mappedB).groupByKey()
# For each (j, list), compute partial products
def compute_products(_, values):
  A_values = []
  B_values = []
```

```
for v in values:
    if v[0] == 'A':
       A values.append((v[1], v[2])) # (i, a ij)
    else:
       B_{values.append((v[1], v[2]))} # (k, b_jk)
  for i, a in A_values:
    for k, b in B_values:
       yield ((i, k), a * b)
products = joined.flatMap(lambda kv: compute_products(kv[0], kv[1]))
# Sum partial results
result = products.reduceByKey(lambda x, y: x + y)
# Display result
for ((i, k), value) in result.collect():
  print(f"C[{i}][{k}] = {value}")
INPUT FILES
1ST FILE SAVE AS matrix_a.txt
A,0,0,1
A,0,1,2
A,1,0,3
A,1,1,4
2<sup>nd</sup> FILE SAVE AS matrix_b.txt
B,0,0,5
B,0,1,6
B,1,0,7
B,1,1,8
```

OUTPUT: cmd:- spark-submit filename.py



3. Develop a Map Reduce program that mines weather data and displays appropriate messages indicating the weather conditions of the day.

```
from pyspark.sql import SparkSession
from pyspark.sql.functions import collect_set
# Step 1: Start Spark
spark = SparkSession.builder.appName("MovieTags").getOrCreate()
# Step 2: Read CSV file
df = spark.read.csv("movietags.csv", header=True)
# Step 3: Select movield and tag columns only
df_tags = df.select("movield", "tag").dropDuplicates()
# Step 4: Group by movield and collect all unique tags
movie_tags = df_tags.groupBy("movieId").agg(collect_set("tag").alias("tags"))
# Step 5: Show result
movie tags.show(truncate=False)
# Step 6: Stop Spark
spark.stop()
INPUT FILE USE BELOW DATASET Save as .csv
userId movieId
                    tag
                           timestamp
      101
             action 1430164910
1
2
      101
             thriller 1430164980
1
      102
             comedy
                           1430164990
3
      101
             action 1430165000
```

OUTPUT

103

2

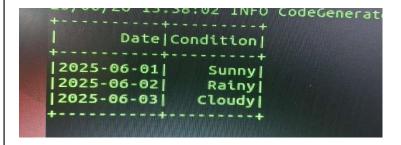


drama 1430165050

4. Develop a MapReduce program to find the tags associated with each movie by analyzing movie lens data.

```
from pyspark.sql import SparkSession
from pyspark.sql.functions import count
# Start Spark
spark = SparkSession.builder.appName("SimpleWeather").getOrCreate()
# Load CSV
df = spark.read.csv("weather.csv", header=True, inferSchema=True)
# Count each condition per day
condition_counts = df.groupBy("Date", "Condition").agg(count("*").alias("count"))
# Sort so most frequent conditions are on top
sorted_counts = condition_counts.orderBy("Date", "count", ascending=[True, False])
# Drop duplicate dates, keeping the first (most frequent)
top_conditions = sorted_counts.dropDuplicates(["Date"])
# Show result
top_conditions.select("Date", "Condition").show()
# Stop Spark
spark.stop()
INPUT FILE DATASET Save as .csv
Date Temperature Condition
2025-06-01
             30
                    Sunny
2025-06-01
             32
                    Sunny
2025-06-01
             31
                    Sunny
2025-06-02
             22
                    Rainy
2025-06-02
             23
                    Cloudy
2025-06-02
             21
                    Rainy
2025-06-03
                    Cloudy
             18
2025-06-03
             19
                    Cloudy
2025-06-03
             20
                    Sunny
```

OUTPUT



5. Implement Functions: Count – Sort – Limit – Skip – Aggregate using MongoDB

```
Create a file:filename.js
const { MongoClient } = require('mongodb');
const uri = 'mongodb://localhost:27017'; // Update if using MongoDB Atlas
const client = new MongoClient(uri);
async function run() {
try {
  await client.connect();
  const db = client.db('testDB');
  const users = db.collection('users');
  // 1. Count
  const count = await users.countDocuments({ city: "New York" });
  console.log("Count of users from New York:", count);
  // 2. Sort (by age descending)
  const sortedUsers = await users.find().sort({ age: -1 }).toArray();
  console.log("Sorted users by age (desc):", sortedUsers);
  // 3. Limit (get top 2)
  const limitedUsers = await users.find().limit(2).toArray();
  console.log("Top 2 users:", limitedUsers);
  // 4. Skip (skip first 2)
  const skippedUsers = await users.find().skip(2).toArray();
  console.log("Users after skipping first 2:", skippedUsers);
```

```
// 5. Aggregate (group by city and count users)
  const aggregation = await users.aggregate([
   { $group: { id: "$city", count: { $sum: 1 } } },
   { $sort: { count: -1 } }
  ]).toArray();
  console.log("User count by city:", aggregation);
 } finally {
  await client.close();
 }
}
run().catch(console.error);
INPUT:
Step 1: Insert Sample Data into MongoDB
>mongosh
Then paste:
use testDB
db.users.insertMany([
 { name: "Alice", age: 25, city: "New York" },
 { name: "Bob", age: 30, city: "Los Angeles" },
 { name: "Charlie", age: 35, city: "Chicago" },
 { name: "David", age: 28, city: "New York" },
 { name: "Eve", age: 22, city: "Chicago" }
])
STEP 2: RUN THE PGM
```

node filename.js

OUTPUT: Total users: 5 Sorted by age: [{ name: 'Charlie', age: 22, ... }, ...] Limit 2 users: [{ name: 'Alice', ... }, { name: 'Bob', ... }] Skip 2 users: [{ name: 'Charlie', ... }, { name: 'David', ... }, ...] Group by city: [{ _id: 'New York', total: 2 }, { _id: 'London', total: 2 }, { _id: 'Paris', total: 1 }]

6. Develop Pig Latin scripts to sort, group, join, project, and filter the data.

INPUT: a. students.txt 1,John,22 2,Alice,23 3,Bob,22 4,David,24 Each line: StudentID, Name, Age b. marks.txt 1,85 2,91 3,78 4,88 Each line: StudentID, Marks **PROGRAM:** -- Load the files A = LOAD 'students.txt' USING PigStorage(',') AS (id:int, name:chararray, age:int); B = LOAD 'marks.txt' USING PigStorage(',') AS (id:int, marks:int); -- Filter students older than 22 C = FILTER A BY age > 22; -- Select only name and age D = FOREACH C GENERATE name, age;

DEPARTMENT OF ARTIFICIAL INTELLIGENCE & DATA SCIENCE **BIG DATA ANALYTICS-Lab Manual**

-- Join student info with marks

E = JOIN A BY id, B BY id;

-- Group by age F = GROUP A BY age; -- Count students in each age group G = FOREACH F GENERATE group AS age, COUNT(A) AS total; -- Sort by age descending H = ORDER A BY age DESC; -- Save output STORE D INTO 'out_filtered' USING PigStorage(','); STORE E INTO 'out_joined' USING PigStorage(','); STORE G INTO 'out_grouped' USING PigStorage(','); STORE H INTO 'out_sorted' USING PigStorage(','); **OUTPUT:** pig -x local student_analysis.pig How to Run It Step 1: Save input files in the current folder ⊕ Copy 'Ø Edit echo -e "1,John,22\n2,Alice,23\n3,Bob,22\n4,David,24" > students.txt echo -e "1,85\n2,91\n3,78\n4,88" > marks.txt Step 2: Save the Pig script ⊕ Copy 10 Edit bash nano student.pig Paste the script above, then save. Step 3: Run Pig in local mode ⊕ Copy * Edit pig -x local student.pig

7. Use Hive to create, alter, and drop databases, tables, views, functions, and indexes.

► How to Run It					
1. Save it as a file, e.g. hive_script.sql.					
2. Open Hive shell:					
bash	∂ Сору	⊘ Edit			
hive					
3. Run the script:					
sql	🗗 Сору	∅ Edit			
SOURCE /full/path/to/hive_script.sql;					
Step 1: Create Database					
CREATE DATABASE mydb;					
Step 2: Use the Database					
USE mydb;					
Step 3: Alter Database Properties					
ALTER DATABASE mydb SET DBPROPERTIES ('edited-by' = 'hive-user');					
Step 4: Create Table					
CREATE TABLE employees (
id INT,					
name STRING,					
salary FLOAT					
)					
ROW FORMAT DELIMITED					
FIELDS TERMINATED BY ','					
STORED AS TEXTFILE:					

Step 5: Alter Table - Add Column
ALTER TABLE employees ADD COLUMNS (age INT);
Step 6: Rename Table
ALTER TABLE employees RENAME TO workers;
Step 7: Create View
CREATE VIEW high_salary AS
SELECT * FROM workers WHERE salary > 5000;
Step 8: Recreate View (simulating ALTER)
DROP VIEW high_salary;
CREATE VIEW high_salary AS
SELECT * FROM workers WHERE salary > 10000;
Step 9: Register Temporary UDF (adjust path/classname if needed)
Step 9: Register Temporary UDF (adjust path/classname if needed) NOTE: This requires the JAR file to exist. Uncomment if you have a UDF.
NOTE: This requires the JAR file to exist. Uncomment if you have a UDF.
NOTE: This requires the JAR file to exist. Uncomment if you have a UDF ADD JAR /path/to/myfunc.jar;
NOTE: This requires the JAR file to exist. Uncomment if you have a UDF ADD JAR /path/to/myfunc.jar;
NOTE: This requires the JAR file to exist. Uncomment if you have a UDF ADD JAR /path/to/myfunc.jar; CREATE TEMPORARY FUNCTION my_upper AS 'com.example.MyUpperFunction';
NOTE: This requires the JAR file to exist. Uncomment if you have a UDF ADD JAR /path/to/myfunc.jar; CREATE TEMPORARY FUNCTION my_upper AS 'com.example.MyUpperFunction'; Step 10: Create Index
NOTE: This requires the JAR file to exist. Uncomment if you have a UDF ADD JAR /path/to/myfunc.jar; CREATE TEMPORARY FUNCTION my_upper AS 'com.example.MyUpperFunction'; Step 10: Create Index CREATE INDEX emp_index ON TABLE workers (id)
NOTE: This requires the JAR file to exist. Uncomment if you have a UDF ADD JAR /path/to/myfunc.jar; CREATE TEMPORARY FUNCTION my_upper AS 'com.example.MyUpperFunction'; Step 10: Create Index CREATE INDEX emp_index ON TABLE workers (id) AS 'COMPACT'
NOTE: This requires the JAR file to exist. Uncomment if you have a UDF ADD JAR /path/to/myfunc.jar; CREATE TEMPORARY FUNCTION my_upper AS 'com.example.MyUpperFunction'; Step 10: Create Index CREATE INDEX emp_index ON TABLE workers (id) AS 'COMPACT'
NOTE: This requires the JAR file to exist. Uncomment if you have a UDF ADD JAR /path/to/myfunc.jar; CREATE TEMPORARY FUNCTION my_upper AS 'com.example.MyUpperFunction'; Step 10: Create Index CREATE INDEX emp_index ON TABLE workers (id) AS 'COMPACT' WITH DEFERRED REBUILD;

Step 12: Drop Index
DROP INDEX emp_index ON workers;
Step 13: Drop View
DROP VIEW high_salary;
Step 14: Drop Table
DROP TABLE workers;
Ston 15: Dron Database
Step 15: Drop Database
DROP DATABASE mydb CASCADE;

8. Implement a word count program in Hadoop and Spark.

INPUT FILE Save as .csv

Virat is talking with ABD

ABD is listening

OUTPUT

```
virat: 1
is: 2
talking: 1
with: 1
ABD: 2
listening: 1
```

FACULTY OF ENGINEERING & TECHNOLOGY (Co-Ed), KALABURAGI SHARNBASVA UNIVERSITY KALABURAGI		
DEPARTMENT OF ARTIFICIAL INTELLIGENCE & BIG DATA ANALYTICS-Lab Manu		